

OMR B9

# Get the code and compile!

- Requirements: git, node.js, and a c++ compiler
  - `sudo apt-get install -y build-essential npm nodejs`
  - Xcode + HomeBrew + `brew install node`

Instructions: Clone the repo, or grab the docker with everything that you need

```
git clone https://github.com/youngar/Base9.git
```

```
cd Base9
```

```
npm install
```

```
Make
```

```
sudo docker run -it jduimovich/b9
```

# What is B9?

- Tiny bytecode interpreter
  - Stack based (opposed to register based)
  - Only supports integers
  - Bytecodes: push\_constant, add subtract, jump, call (See b9.h)
- Project contains:
  - b9 – bytecode interpreter
  - b9.js – compiles language to bytecode
- Compilation process to bytecodes:
  - program.src → myprogram.cpp → myprogram.so
- To run: `./b9 myprogram.so`

# Example: program.src → program.cpp

```
Function addOneAndTwo() { return 1+2 }
```

## Compiles to:

```
Instruction addOneAndTwo[] = {  
    decl (0,8),           // (args,temps)  assume max 8 temps  
    decl (0,0),           // 0: space for JIT address  
    decl (0,0),           // 1: space for JIT address  
    createInstruction(PUSH_CONSTANT,1), // number constant 1  
    createInstruction(PUSH_CONSTANT,2), // number constant 2  
    createInstruction(ADD,0),  
    createInstruction(RETURN,0), // return  
    createInstruction(NO_MORE_BYTECODES, 0)};
```

# Try it:

- Edit program.src
- Run `make program`

This will:

1. Compile the program
2. Print produced bytecodes from program.cpp
3. Run the program with just the interpreter
4. Run the program by compiling all functions

```
./b9 program.so
```

```
Loading "program.so
```

```
"Handle=0x1cc6de0 table=0x7f43d489a060
```

```
Running Interpreted
```

```
result is 3
```

```
Running JIT
```

```
generating index=3 bc=PUSH_CONSTANT(1) param=1
```

```
generating index=4 bc=PUSH_CONSTANT(1) param=2
```

```
generating index=5 bc=ADD(6) param=0
```

```
generating index=6 bc=RETURN(8) param=0
```

```
Compiled success address = <0x7f43d489b034>
```

```
Done gen code
```

```
result is 3
```

# How bytecode interpreters work

- Bytecodes
- Stack
- Main interpreter loop
- How JIT callouts work

# Lets add new bytecodes MUL and DIV

- Adding two new bytecodes,
  - Multiply -- MUL, pops two elements off the stack, push the product
  - Divide – DIV, pops two elements of the stack, pushes the quotient
- The compiler, b9.js, already supports for divide and multiply:

```
function multiply() {  
    return 5*6  
}
```

```
createInstruction(PUSH_CONSTANT, 5), // number constant 5  
createInstruction(PUSH_CONSTANT, 6), // number constant 6  
createInstruction(MUL, 0),
```

# Our test program: program.src

```
function main_function() {  
    var a = 1*2;  
    return multiply_by_4(a*3);  
}
```

```
function multiply_by_4(a) {  
    return a*4;  
}
```

- Expected output is 24



# Teach the interpreter MUL

In b9.h:

1. Need to disable (comment out) 3 JIT features for now:

```
// #define USE_DIRECT_CALL    0  
// #define PASS_PARAMETERS_DIRECTLY  0  
// #define USE_VM_OPERAND_STACK  0
```

2. Add the bytecode value for MUL

```
#define MUL 0xB
```

# Teach the interpreter MUL

- Add a function to perform the multiplication:

```
Void bc_mul(ExecutionContext *context) {  
    // a*b is push(a);push(b); multiply  
    StackElement right = pop(context);  
    StackElement left = pop(context);  
    StackElement result = left * right;  
    push(context, result);  
}
```

# Teach the interpreter MUL

- In main.cpp, function:

```
interpret (ExecutionContext *context, Instruction
*program)
```

- Add a handler for MUL:

```
while (*instructionPointer != NO_MORE_BYTECODES) {
. . .
    Case MUL:
        bc_mul(context);
        break;
. . .
```

# Test it!

- Edit program.src to use multiplication, e.g.  $a * b$ 
  - Run ``make program``
- Edit test.cpp to use the MUL opcode
  - Run ``make test``
- The interpreter times interpreted mode vs JITed mode
  - But the JIT doesn't know how to compile MUL

# Teach the JIT MUL

- In b9jit.cpp, add a handler to generate IL

```
B9Method::handle_bc_mul(TR::BytecodeBuilder
*builder,
    TR::BytecodeBuilder *nextBuilder) {
    TR::IlValue *right = pop(builder);
    TR::IlValue *left = pop(builder);
    push(builder, builder->Mul(left, right));
    builder->AddFallThroughBuilder(nextBuilder);
}
```

# Teach the JIT MUL

- In `b9jit.cpp`, function:

`B9Method::generateILForBytecode`

- Add a handler to generate IL

```
while (*instructionPointer != NO_MORE_BYTECODES) {  
    . . .  
    case MUL:  
        handle_bc_mul(builder, nextBytecodeBuilder);  
        break;  
    . . .  
}
```

# Test it

- Edit program.src to use multiplication, e.g.  $a * b$ 
  - Run ``make program``
- Edit test.cpp to use the MUL opcode
  - Run ``make test``
- Verify your interpreter and JIT are getting the same result!
- What is the speed up?

# Test it

```
function program() {  
    var a = 1*2;  
    var b = 3*4;  
    var c = 5*6;  
    return a*b*c;  
}
```

```
./b9 program.so -loop 1000000
```

Result for Interp is 720, resultJit is 720

Time for Interp 471 ms, JIT 11 ms

JIT speedup = **42.818182**



# Quickly Adding JIT to pre-existing Languages

- Take your already existing BC implementation:

```
Void bc_mul (ExecutionContext *context;
```

- And JIT version just generates calls it:

```
handle_bc_mul (...) {  
    builder->Call ("bc_mul", 2,  
        builder->Load ("context")) ;  
}
```

# Conclusions

- What is OMR used for?
- Who is OMR For?
- As student, why should you be interested?



# Our test program:

```
function main_function() {  
    var a = factorial(1000)  
    return 0;  
}  
  
function factorial(a) {  
    if (a == 0)  
        return 1;  
  
    return a * factorial(a - 1)  
}
```

# Making it faster!

- The generated code for MUL always pushes and pops to the Interpreter's stack
- What if we wanted to keep the values in registers, and omit the pushes and pops to the stack?
- Problem: When we call another function, we enter the interpreter again
  - The interpreter's stack must be correct and up to date, so how can we omit stack pushes and pops?

# Introducing VirtualMachineState

- Answer: simulate the interpreter's stack while in a compiled method, and restore the correct stack when calling back into the interpreter
- Call Commit() before a CALL
  - Store values from the simulated stack into the interpreter's stack
- Call Reload() after a CALL
  - Reloads values from the interpreter's stack into the simulated stack
- When compiling PUSH and POP, just simulate the calls
- Base9 codebase has this implemented, you can enable this by uncommenting in B9.h

```
#define USE_VM_OPERAND_STACK 1
```

# Introducing VirtualMachineState

```
B9Method::pop(TR::BytecodeBuilder *builder) {  
    // simulate a pop  
    builder->vmState()->_stack->Pop(builder);  
}
```

```
B9Method::push(TR::BytecodeBuilder *builder,  
               TR::IlValue *value) {  
    // simulate a push  
    builder->vmState()->_stack->Push(builder, value);  
}
```

# Introducing VirtualMachineState

- Make sure the stack restored before calling back to the interpreter

case CALL:

```
builder->vmState() ->Commit(builder)
```

```
result = builder->Call("interpret", . . .);
```

```
break;
```

- No call to Reload() after calling to the interpreter?
  - Our interpreter returns the stack to exactly how it was before



# Even faster?

- Generated code for CALL always goes into the interpreter
  - Interpreter checks to see if a method was compiled
- Solution: We can check if we're trying to CALL a compiled method at compile time
- Pseudo code for generating a call:

```
If (functionToCall == currentFunction ||
    functionToCall.hasBeenJitted) {
    generateCallTo(functionToCall);
} else
    generateCallTo(interpreter);
}
```

# Even faster?

- B9 will compile methods in the order they are declared..
  - Need to make sure we JIT things in the correct order..
    - Make sure your program.src has its functions in the right order!
- Base9 codebase has this implemented, you can enable this by uncommenting in B9.h

```
#define USE_DIRECT_CALL 1
```

# Even more faster?

- What if it calls into compiled code was calling compiled code?
  - Don't have to push arguments on to the stack
  - Don't have to Commit() and Restore() the stack
  - Our benchmark is a recursive function..

- Psudo Code

```
if (functionToCall.isCompiled()) {  
    /* pop from the simulated stack */  
    TR::IlValue * p1 = pop(builder);  
    result = builder->Call(functionToCall, argC, p1);  
    /* Return values still on the stack, commit it to to the simulated stack */  
} else {  
  
}
```

# Getting it working

- Loading function arguments needs to use the simulated stack
  - Local variables are on the same stack, they can use the simulated stack as well
- The interpreter needs to call compiled code using C calling convention
  - i.e. pop parameters off the it's stack, and pass them through a C call
  - Old calling convention:
    - `push(1); push(2); result = *compiledFunction(stack, stackPointer);`
  - New calling convention:
    - `result = *compiledFunction(1, 2);`